

# How Many In-Context Examples Is a Task Description Worth?

Bruce Quan Nguyen

[Affiliation]

Draft of 30 September 2026

## Abstract

While prompts in modern language models typically consist of both a task description and in-context examples, theoretical frameworks for in-context learning (ICL) predominantly focus on the latter. In this work, we adopt a Bayesian perspective on ICL to formally quantify the value of task descriptions. By casting in-context examples as observational data and the task description as a prior, we can measure the description’s value via its prior effective sample size (ESS)—the equivalent number of examples it provides. We analyze this in the setting of in-context linear regression, where Bayes-optimal predictions can be computed exactly. We define the ESS of a description by matching the one-step log-loss regret of a Bayes-optimal predictor using only the description to one using only examples. Our analysis reveals that a description’s worth is governed by two factors: its precision (how much it reduces prior variance) and its reliability (the probability it is correct). We show that for a reliable description, removing a fraction  $1 - r$  of the prior variance in a  $d$ -dimensional task yields an ESS of approximately  $d(1 - r)$ . However, for unreliable descriptions, the ESS saturates: the inherent uncertainty imposes a regret floor, so that beyond a certain point increasing the precision of the description yields diminishing returns: at  $p = 0.9$ , a fiftyfold increase in precision adds a few examples where a reliable description gains a hundred. Finally, small meta-trained Transformers reproduce these ESS values when trained on the matching distribution, and a Transformer trained only on reliable descriptions treats an unreliable one as a Bayes-optimal predictor with full trust would: a precise description becomes worth almost nothing. Our claims concern Bayes-optimal predictors and small meta-trained sequence models, not large language models.

## 1 Introduction

In-context learning (ICL) allows models to adapt to new tasks from a prompt (Brown et al., 2020). In practice, a prompt generally comprises a natural language task description paired with a few input-output demonstrations. However, formal theories of ICL largely abstract prompts as mere collections of demonstrations (Garg et al., 2022; Akyürek et al., 2023; Xie et al., 2022; Genewein et al., 2025). This leaves a significant gap in our understanding, as empirical evidence on the value of task descriptions remains mixed. Some studies show that a good description can be worth hundreds of data points (Le Scao and Rush, 2021; Honda et al., 2025), whereas others find that misleading descriptions perform about as well as informative ones (Webson and Pavlick, 2022) and that enough examples override an explicitly stated bias (Gupta et al., 2025). Even recent theoretical models that incorporate descriptors (Huang and Ge, 2025; Lin et al., 2025; Tong et al., 2026) fail to express the utility of these descriptions on a commensurate scale with the examples themselves.

We quantify the value of a task description as the number of in-context examples it can replace. Following the Bayesian interpretation of ICL (Xie et al., 2022), a sequence model can be seen as performing posterior predictive inference over a latent task. In this view, in-context examples are observations, and the task description acts as a prior. Within Bayesian statistics, the weight of a prior is naturally measured by its *prior effective sample size* (ESS) (Morita et al., 2008). This ESS is primarily driven by two dimensions: the prior’s concentration, and its concordance with the underlying truth (Evans and Moshonov, 2006). Translating this to task descriptions, we identify two key characteristics: **precision** (the extent to which the description reduces posterior uncertainty) and **reliability** (the probability that the description is actually correct).

We compute the ESS as a function of precision and reliability within in-context linear regression (Garg et al., 2022), a setting where the Bayes-optimal predictor is exact. We focus on one-step prediction, mirroring standard prompt usage.

**Contributions.**




Figure 1: **A task description is worth a number of in-context examples, set by its precision and capped by its reliability.** (a) We read a prompt as Bayesian inference over a latent task  $w$ : the description sets the prior and the examples are data. The description’s effective sample size (ESS) is the number of examples  $k$  that lowers a Bayes-optimal predictor’s one-step regret as much as the description does. (b) Three descriptions and their ESS. Precision  $1/r$  sets how narrowly the prior concentrates around the stated  $m$ ; reliability  $p$  is the chance the description is correct, and otherwise  $w$  follows the base prior (dashed; its share  $1 - p$  shaded pink). Prior shapes are sketches; ESS values are from Table 2 ( $d = 16, a_0 = 10$ ).

1. We introduce a formal definition for the ESS of a task description by equating the one-step log-loss regret of a description to that of  $n$  examples (§3).
2. We establish that for a perfectly reliable description, the ESS scales linearly as  $\text{ESS} \approx d(1 - r)$ , where  $d$  is the dimensionality and  $r$  is the precision ratio (§4).
3. We demonstrate that unreliable descriptions induce a hard floor on regret, governed by the description’s error probability  $1 - p$  and its entropy  $H(p)$ . Consequently, the ESS saturates; making an unreliable description infinitely precise yields diminishing returns (§5).
4. We meta-train small Transformers on this process and show that they reproduce the Bayes-optimal ESS when trained on the matching distribution, and that a model trained on reliable descriptions does not discount an unreliable one (§6).

## 2 Related Work

**Bayesian Perspectives on ICL.** ICL can be framed as implicit Bayesian inference (Xie et al., 2022). In the context of linear regression, meta-trained Transformers effectively implement Bayes-optimal ridge regression (Akyürek et al., 2023; Raventós et al., 2023; Panwar et al., 2024). Genewein et al. (2025) evaluate such models using log-loss regret against the Bayes-optimal baseline, a metric we similarly adopt. Yet, these works do not translate the influence of a prior into an equivalent number of in-context examples.

**Task Descriptions in ICL Theory.** Recent studies have begun exploring task descriptions theoretically. Huang and Ge (2025) include a descriptor token representing the input mean, but because it is independent of the regression weights, its ESS in our framework would be zero. Lin et al. (2025) guide models using hypothesis classes, and Tong et al. (2026) demonstrate how models can exponentially override incorrect instructions given enough examples. Lin and Lee (2024) analyze Gaussian-mixture priors similar to ours, but none of these works quantify the description’s value in terms of sample size.

**Effective Sample Size.** In statistics, Morita et al. (2008) define ESS via curvature matching, while Reimherr et al. (2021) highlight the need for prediction-based ESS measures. Our approach directly addresses this by defining ESS through predictive regret. Furthermore, our model of reliability draws on robust mixture priors (Schmidli et al., 2014), commonly used in clinical trials to incorporate potentially conflicting historical data.

### 3 Problem Formulation

**Task Distribution.** We consider tasks defined by a weight vector  $w \in \mathbb{R}^d$ . An in-context example consists of an input  $x \sim \mathcal{N}(0, I_d)$  and a corresponding label  $y = w^\top x + \epsilon$ , where the noise is  $\epsilon \sim \mathcal{N}(0, \sigma_y^2)$ . The base prior over tasks is  $w \sim \mathcal{N}(0, s_0^2 I_d)$ . For simplicity, we set  $\sigma_y = 1$  and define the prior signal-to-noise ratio (SNR) as  $a_0 = s_0^2 / \sigma_y^2$ .

**Task Description Model.** We model a task description as a vector  $m \in \mathbb{R}^d$  that provides information about  $w$ . The description’s *precision ratio* is  $r \in (0, 1)$ , meaning the variance around the description is  $s_\ell^2 = r s_0^2$ . The description is generated hierarchically via  $m \sim \mathcal{N}(0, (1 - r) s_0^2 I_d)$ , ensuring the marginal distribution of  $w$  matches the base prior. A description is not always correct. We define its *reliability*  $p$  as the probability that the description is accurate. If correct,  $w \sim \mathcal{N}(m, s_\ell^2 I_d)$ ; if incorrect (with probability  $1 - p$ ),  $w$  simply follows the base prior  $\mathcal{N}(0, s_0^2 I_d)$  independently of  $m$ . Consequently, the Bayes-optimal prior conditioned on the description is a robust mixture (Schmidli et al., 2014):  $p \mathcal{N}(m, s_\ell^2 I_d) + (1 - p) \mathcal{N}(0, s_0^2 I_d)$ . This formulation treats the description as a weighted numerical prior rather than raw text to be parsed.

**Regret and Effective Sample Size.** We evaluate predictors based on one-step predictive log-loss for a new query  $x$ . The regret is the expected difference in log-loss between the predictor and an oracle that knows  $w$  exactly, the expectation taken over tasks, inputs, and labels (Genewein et al., 2025).

**Definition 1** (Effective Sample Size). *Let  $R^{\text{desc}}$  be the regret of a Bayes-optimal predictor given only the task description, and  $R^{\text{ex}}(n)$  be the regret of a Bayes-optimal predictor under the base prior given  $n$  examples. The effective sample size (ESS) of the description is the value  $n$  (interpolated continuously) where  $R^{\text{ex}}(n) = R^{\text{desc}}$ .*

Since we compare Bayes-optimal predictors under different conditioning sets, our ESS fundamentally measures information content rather than the specific architecture of a learner.

### 4 The Value of Precision

We first analyze a perfectly reliable description ( $p = 1$ ). Given  $n$  examples  $X_n \in \mathbb{R}^{n \times d}$ , the posterior covariance under a Gaussian prior with SNR  $a$  is  $\Sigma = \sigma_y^2 (I/a + X_n^\top X_n)^{-1}$ . This is independent of the examples’ labels and the description  $m$ . The expected one-step regret simplifies to  $\frac{1}{2} \mathbb{E} \log(1 + x^\top \Sigma x)$ . A reliable description equates to  $n = 0$  and  $a = r a_0$ , whereas relying solely on examples uses  $a = a_0$ .

**Proposition 1.** *Let  $\psi$  be the digamma function. As  $a_0 \rightarrow \infty$ , for  $n < d$ ,*

$$R^{\text{ex}}(n) = \frac{1}{2} \log a_0 + \frac{1}{2} \log 2 + \frac{1}{2} \psi\left(\frac{d-n}{2}\right) + o(1), \quad R^{\text{desc}} = \frac{1}{2} \log(r a_0) + \frac{1}{2} \log 2 + \frac{1}{2} \psi\left(\frac{d}{2}\right) + o(1),$$

*so the high-SNR ESS, with  $R^{\text{ex}}(n)$  extended to non-integer  $n$  through the digamma function, solves  $\psi\left(\frac{d-n}{2}\right) - \psi\left(\frac{d}{2}\right) = \log r$ . As  $rd \rightarrow \infty$ , uniformly in  $r \in (0, 1)$ , its solution is*

$$\text{ESS} = (d - 1)(1 - r) + O\left(\frac{1}{rd}\right).$$

Table 1: Exact ESS of a reliable description ( $p = 1$ ) for  $a_0 = 10$ , averaged over three seeds (standard deviation across seeds below 2% of each value). The approximation  $d(1 - r)$  holds well as long as the exact ESS does not exceed  $d$ .

| $d$ | $r$  | ESS  | $d(1 - r)$ |
|-----|------|------|------------|
| 4   | 0.5  | 1.77 | 2.0        |
| 4   | 0.2  | 3.32 | 3.2        |
| 16  | 0.5  | 7.62 | 8.0        |
| 16  | 0.2  | 12.6 | 12.8       |
| 16  | 0.05 | 16.9 | 15.2       |
| 16  | 0.01 | 26.1 | 15.8       |
| 64  | 0.5  | 31.6 | 32.0       |
| 64  | 0.2  | 50.7 | 51.2       |
| 64  | 0.05 | 62.1 | 60.8       |
| 64  | 0.01 | 73.5 | 63.4       |

*Idea.* Regret is half the log of the leftover variance in the prediction. At high SNR each example removes all the variance along the one direction it looks at, so  $n < d$  examples leave  $d - n$  directions untouched; a description shrinks all  $d$  directions by the factor  $r$ . Matching the two leftovers gives  $d - n \approx rd$ . The proof, with the error term, is in Appendix A.

The error term makes the range of validity precise: the approximation needs  $rd \gg 1$ , that is, the description must leave many more than one example’s worth of variance unexplained, which is the condition  $\text{ESS} < d$  in the first-order form  $\text{ESS} \approx d(1 - r)$ . At  $a_0 = 100$  the exact ESS is within 0.1 of  $(d - 1)(1 - r)$  for  $d \in \{16, 64\}$  and  $r \in \{0.5, 0.2\}$  (values 7.50, 12.1, 31.6, 50.5 against 7.5, 12.0, 31.5, 50.4).

Table 1 illustrates this alignment. A description eliminating 50% of the prior variance is worth roughly half the dimension in examples. However, once the description is precise enough that its ESS surpasses  $d$ , the approximation breaks down. The true ESS scales beyond  $d$  because  $d$  noisy examples cannot pinpoint  $w$  perfectly.

## 5 The Bottleneck of Reliability

In reality, descriptions can be flawed or ambiguous. When providing an unreliable description without examples, the Bayes-optimal predictive distribution becomes a mixture. Without observations that could identify the correct component, the predictor must average over both.

**Proposition 2.** Let  $R^{\text{desc}}(p, r)$  be the regret of a description with reliability  $p$  and precision ratio  $r$ ,  $R_{p=1}^{\text{desc}}(r)$  the same for a reliable one, and  $R^{\text{ex}}(0)$  the regret with neither description nor examples. For every  $p \in (0, 1)$  and  $r \in (0, 1)$ ,

$$p R_{p=1}^{\text{desc}}(r) + (1 - p) R^{\text{ex}}(0) \leq R^{\text{desc}}(p, r) \leq p R_{p=1}^{\text{desc}}(r) + (1 - p) R^{\text{ex}}(0) + H(p),$$

where  $H(p)$  is the binary entropy in nats. In particular  $R^{\text{desc}}(p, r) \geq (1 - p) R^{\text{ex}}(0)$  for every  $r$ , and as  $r \rightarrow 0$  the regret lies between  $(1 - p) R^{\text{ex}}(0)$  and  $(1 - p) R^{\text{ex}}(0) + H(p)$ .

*Idea.* Compare the learner with a twin who is also told whether the description is correct. The twin’s regret is  $p R_{p=1}^{\text{desc}}(r) + (1 - p) R^{\text{ex}}(0)$ . The learner cannot beat the twin, since extra information never hurts a Bayes-optimal predictor. And it loses at most  $\log(1/p)$  nats when the description is correct and  $\log(1/(1 - p))$  when it is wrong, because its predictive puts those weights on the twin’s answer; the average is  $H(p)$ . The proof is in Appendix B.

The gap between the two bounds is the cost of not knowing whether the description is correct, and it is at most  $H(p)$  however precise the description. For  $d = 16$  and  $a_0 = 10$ ,  $R^{\text{ex}}(0) = 2.51$  nats and  $R_{p=1}^{\text{desc}}(0.001) = 0.07$ ; at  $p = 0.9$  the two bounds are 0.32 and 0.64, and the computed regret at  $r = 0.001$  is 0.57.

Because  $R^{\text{ex}}(n)$  decreases with  $n$ , the lower bound caps the ESS at the  $n$  where  $R^{\text{ex}}(n) = (1 - p) R^{\text{ex}}(0)$ , however precise the description. An unreliable description can replace only a bounded number of examples, and the identification cost, up to  $H(p)$ , lowers the cap further.

Table 2 demonstrates this saturation. For a perfectly reliable description, refining  $r$  from 0.05 to 0.001 raises the ESS from 16.9 to 116.5. However, if  $p = 0.9$ , the ESS only grows from 14.8 to 23.1. Even a 1% unreliability ( $p = 0.99$ ) nearly halves the maximum potential ESS of the most precise description. The same saturation patterns manifest at  $d \in \{4, 64\}$  and  $a_0 = 100$ .

Table 2: ESS of unreliable descriptions ( $d = 16$ ,  $a_0 = 10$ ), three seeds; standard deviation across seeds at most 2% of each value. As precision increases (smaller  $r$ ), the ESS is capped by the reliability  $p$ .

| $r$   | $p = 1$ | $p = 0.99$ | $p = 0.9$ | $p = 0.7$ |
|-------|---------|------------|-----------|-----------|
| 0.5   | 7.6     | 7.5        | 6.3       | 4.1       |
| 0.05  | 16.9    | 16.6       | 14.8      | 11.7      |
| 0.001 | 116.5   | 64.2       | 23.1      | 15.6      |

Table 3: Networks against the Bayes-optimal predictor ( $d = 5$ ,  $a_0 = 10$ , one trained model per row, 20,000 evaluation prompts). “Trained” is the Bayes-optimal predictor with trust equal to the training reliability  $p_{\text{train}}$ , the predictor each network is trained toward; “calibrated” is the predictor with trust equal to the test reliability, shown where the two differ. The last two columns are the mean regret gap, network minus Bayes-optimal, over  $n = 0, \dots, 15$ , with and without a descriptor.

| $r$  | $p_{\text{train}}$ | $p_{\text{test}}$ | ESS     |         |            | Regret gap (nats) |       |
|------|--------------------|-------------------|---------|---------|------------|-------------------|-------|
|      |                    |                   | Network | Trained | Calibrated | Descriptor        | None  |
| 0.05 | 1                  | 1                 | 6.7     | 6.7     | —          | 0.023             | 0.022 |
| 0.5  | 1                  | 1                 | 2.2     | 2.3     | —          | 0.019             | 0.019 |
| 0.05 | 0.9                | 0.9               | 3.9     | 5.0     | —          | 0.070             | 0.025 |
| 0.5  | 0.9                | 0.9               | 1.8     | 1.9     | —          | 0.019             | 0.016 |
| 0.05 | 1                  | 0.9               | 0.7     | 0.0     | 5.0        | -0.000            | 0.022 |
| 0.5  | 1                  | 0.9               | 1.6     | 1.7     | 1.9        | 0.024             | 0.019 |

## 6 Meta-Trained Transformers

**Design.** To test whether a learned predictor reproduces these values, we meta-train Transformers on prompts generated by the process of §3 and compare them with the Bayes-optimal predictor on the same prompts, following Genewein et al. (2025). Prompts use the prefix embedding of Huang and Ge (2025): an optional descriptor token carrying  $m$ , up to 15 example tokens each holding an input beside its answer, and a final query token, with two indicator coordinates marking the token type. There is no positional encoding, so the examples are exchangeable. The architecture and optimisation follow Garg et al. (2022): 12 layers, 8 heads, width 256, Adam at  $10^{-4}$ , fresh prompts at every step. The network outputs a mean and a log variance for the answer to the query and is trained on the Gaussian log loss, so its regret is on the scale of §3. We set  $d = 5$  and  $a_0 = 10$ , draw the number of examples uniformly from 0 to 15, and omit the descriptor in half the prompts. The descriptor supplies only  $m$ : each network is trained at one precision  $r$  and one reliability  $p$ , which it never sees and must learn from the training distribution. Every model trains for 40,000 steps at batch 1024, one model per setting. A pilot confirmed the budget: the mean gap between the network and the Bayes-optimal predictor falls from 0.07 nats at 5,000 steps to 0.02 at 20,000 and stays there to 40,000.

**Readouts.** We evaluate each network on 20,000 fresh prompts for every  $n$  from 0 to 15, with and without a descriptor, and compute its ESS exactly as for the Bayes-optimal predictor: the regret with the descriptor and no examples is placed on the network’s own examples-only regret curve. We also evaluate the two models trained at  $p = 1$  on prompts whose descriptors are correct only with probability 0.9, which asks how a learner that has never seen a wrong description treats one.

**Results.** Table 3 and Figure 2 give the readouts. For reliable descriptors the networks reproduce the Bayes-optimal values: ESS 6.7 against 6.7 for the precise descriptor and 2.2 against 2.3 for the coarse one, with regret within 0.02 nats of the Bayes-optimal predictor at every  $n$ . The coarse unreliable descriptor is also reproduced, 1.8 against 1.9. The precise unreliable descriptor is not: the network’s ESS is 3.9 against 5.0, its regret from the descriptor alone is 0.3 nats above the Bayes-optimal predictor’s, and the trust that best explains its predictions is 0.85 rather than 0.9. The gap closes as examples arrive (Figure 2c). *[The training loss of this model was still falling slowly at 40,000 steps; a continuation to 80,000 steps is running to separate under-training from under-trust. Update this paragraph and the table row from its result.]*




Figure 2: **Networks trained on the matching distribution track the Bayes-optimal predictor, except the precise unreliable descriptor.** Regret against the number of examples for the network (markers, standard errors over 20,000 prompts) and the Bayes-optimal predictor (lines), with and without a descriptor; one trained model per panel. (a)  $r = 0.05, p = 1$ ; (b)  $r = 0.5, p = 1$ ; (c)  $r = 0.05, p = 0.9$ ; (d)  $r = 0.5, p = 0.9$ . In (c) the network’s regret with the descriptor alone exceeds the Bayes-optimal predictor’s by 0.3 nats, and its ESS is 3.9 against 5.0.

**Transfer to unreliable descriptors.** The models trained at  $p = 1$  behave like a Bayes-optimal predictor that trusts the descriptor fully. On prompts where one descriptor in ten is wrong, the model trained on precise reliable descriptors gets almost nothing from a descriptor: ESS 0.7, where a calibrated predictor would get 5.0, and the fully trusting Bayes-optimal predictor gets 0. Its implied trust is 0.99. The coarse model is forgiven its trust, 1.6 against the calibrated 1.9, because a coarse descriptor that is wrong does little harm. This is the mechanism of §5 seen in a trained learner: the cost of unreliability falls on precise descriptors, and a learner that has never seen a wrong descriptor pays it in full.

## 7 Discussion

Our analysis offers a precise vocabulary for discussing the utility of task descriptions in prompts. For a Bayes-optimal predictor, a task description is worth a specific number of in-context examples. Assuming perfect reliability, this worth scales linearly with the proportion of prior variance the description eliminates. However, reliability acts as a strict bottleneck. Because a predictor must hedge against the possibility that the description is incorrect, the value of very precise descriptions saturates if they are not perfectly reliable.

**Limitations and Future Work.** Our theoretical claims are restricted to Bayes-optimal predictors and small-scale meta-trained Transformers in a linear-Gaussian setting, chosen to permit exact calculation. Scaling these insights to large language models operating on complex text remains an open challenge. Additionally, our ESS metric applies to one-step prediction. Over long horizons, sequential feedback allows predictors to detect and adapt to incorrect descriptions, meaning ESS will dynamically evolve with the prediction horizon—a compelling avenue for future research.

## References

- Ekin Akyürek, Dale Schuurmans, Jacob Andreas, Tengyu Ma, and Denny Zhou. What learning algorithm is in-context learning? Investigations with linear models. In *ICLR*, 2023.
- Tom B. Brown et al. Language models are few-shot learners. In *NeurIPS*, 2020.
- Michael Evans and Hadas Moshonov. Checking for prior-data conflict. *Bayesian Analysis*, 1(4):893–914, 2006.
- Shivam Garg, Dimitris Tsipras, Percy Liang, and Gregory Valiant. What can transformers learn in-context? A case study of simple function classes. In *NeurIPS*, 2022.
- Tim Genewein, Li Kevin Wenliang, Jordi Grau-Moya, Anian Ruoss, Laurent Orseau, and Marcus Hutter. Understanding prompt tuning and in-context learning via meta-learning. In *NeurIPS*, 2025.
- Ritwik Gupta, Rodolfo Corona, Jiaxin Ge, Eric Wang, Dan Klein, Trevor Darrell, and David M. Chan. Enough coin flips can make LLMs act Bayesian. In *ACL*, 2025.
- Ukyo Honda, Soichiro Murakami, and Peinan Zhang. Distilling many-shot in-context learning into a cheat sheet. In *Findings of EMNLP*, 2025.
- Ruomin Huang and Rong Ge. Task descriptors help transformers learn linear models in-context. In *ICLR*, 2025.
- Teven Le Scao and Alexander M. Rush. How many data points is a prompt worth? In *NAACL*, 2021.
- Ziqian Lin and Kangwook Lee. Dual operating modes of in-context learning. In *ICML*, 2024.
- Ziqian Lin, Shubham Kumar Bharti, and Kangwook Lee. In-context learning with hypothesis-class guidance. *arXiv:2502.19787*, 2025.
- Satoshi Morita, Peter F. Thall, and Peter Müller. Determining the effective sample size of a parametric prior. *Biometrics*, 64(2):595–602, 2008.
- Madhur Panwar, Kabir Ahuja, and Navin Goyal. In-context learning through the Bayesian prism. In *ICLR*, 2024.
- Allan Raventós, Mansheej Paul, Feng Chen, and Surya Ganguli. Pretraining task diversity and the emergence of non-Bayesian in-context learning for regression. In *NeurIPS*, 2023.
- Matthew Reimherr, Xiao-Li Meng, and Dan L. Nicolae. Prior sample size extensions for assessing prior impact and prior-likelihood discordance. *Journal of the Royal Statistical Society: Series B*, 83(3):413–437, 2021.
- Heinz Schmidli, Sandro Gsteiger, Satrajit Roychoudhury, Anthony O’Hagan, David Spiegelhalter, and Beat Neuen-schwander. Robust meta-analytic-predictive priors in clinical trials with historical control information. *Biometrics*, 70(4):1023–1032, 2014.
- X. Tong, Y. Zeng, and J. Zhang. Demonstrations, CoT, and prompting: a theoretical analysis of ICL. *arXiv:2603.19611*, 2026.
- Albert Webson and Ellie Pavlick. Do prompt-based models really understand the meaning of their prompts? In *NAACL*, 2022.
- Sang Michael Xie, Aditi Raghunathan, Percy Liang, and Tengyu Ma. An explanation of in-context learning as implicit Bayesian inference. In *ICLR*, 2022.

## A Proof of Proposition 1

*Idea.* Regret is half the log of the leftover variance in the prediction. At high SNR each example removes all the variance along the one direction it looks at and none elsewhere, so  $n < d$  examples leave  $d - n$  directions untouched. A description instead shrinks every one of the  $d$  directions by the factor  $r$ . The ESS is the  $n$  at which the two leftovers match.

*Proof. Step 1: regret is leftover variance.* The Bayes predictive for  $y$  is the true conditional law of  $y$  given the data, a Gaussian with variance  $1 + x^\top \Sigma x$  where  $\Sigma$  is the posterior covariance of  $w$ . The expected log loss of a true conditional law is its entropy,  $\frac{1}{2} \log(2\pi e(1 + x^\top \Sigma x))$ ; the oracle's predictive is  $\mathcal{N}(w^\top x, 1)$  with entropy  $\frac{1}{2} \log(2\pi e)$ . So the regret is  $\frac{1}{2} \mathbb{E} \log(1 + x^\top \Sigma x)$ : half the log of the leftover variance.

*Step 2: examples.* Write  $X_n^\top X_n = \sum_{i \leq n} \lambda_i u_i u_i^\top$  with  $\lambda_i > 0$  almost surely, and let  $P_\perp$  project onto the  $d - n$  directions the examples did not look at. Then  $\Sigma = (I/a_0 + X_n^\top X_n)^{-1}$  has eigenvalue  $a_0$  on those directions and  $1/(1/a_0 + \lambda_i)$  on the others, so

$$1 + x^\top \Sigma x = a_0 \|P_\perp x\|^2 + T, \quad T = 1 + \sum_{i \leq n} \frac{(u_i^\top x)^2}{1/a_0 + \lambda_i}.$$

Because  $x$  is a fresh isotropic Gaussian,  $\|P_\perp x\|^2 \sim \chi_{d-n}^2$ , independent of  $X_n$ . Hence

$$R^{\text{ex}}(n) = \frac{1}{2} \log a_0 + \frac{1}{2} \mathbb{E} \log \chi_{d-n}^2 + \frac{1}{2} \mathbb{E} \log \left( 1 + \frac{T/a_0}{\|P_\perp x\|^2} \right).$$

The last term is nonnegative, and  $T/a_0 = 1/a_0 + \sum_{i \leq n} (u_i^\top x)^2 / (1 + a_0 \lambda_i)$  decreases to 0 as  $a_0 \rightarrow \infty$ , so by monotone convergence the term is  $o(1)$ . (It is finite at any finite  $a_0$  because  $R^{\text{ex}}(n)$  is.)

*Step 3: description.* Here  $n = 0$  and  $\Sigma = r a_0 I$ , so  $1 + x^\top \Sigma x = 1 + r a_0 \|x\|^2$  with  $\|x\|^2 \sim \chi_d^2$ , and the same argument gives  $R^{\text{desc}} = \frac{1}{2} \log(r a_0) + \frac{1}{2} \mathbb{E} \log \chi_d^2 + o(1)$ .

*Step 4: match them.* The  $\log a_0$  terms cancel, and  $\mathbb{E} \log \chi_k^2 = \psi(k/2) + \log 2$ . Taking  $\psi((d-n)/2)$  as the extension of  $R^{\text{ex}}(n)$  to non-integer  $n$ , the high-SNR ESS solves

$$\psi\left(\frac{d-n}{2}\right) - \psi\left(\frac{d}{2}\right) = \log r,$$

which has one root in  $(0, d)$  because  $\psi$  is increasing and  $\log r < 0$ . Read it as:  $d - n$  untouched directions must be as big, on a log scale, as  $d$  directions each shrunk by  $r$ . Replacing  $\mathbb{E} \log \chi_k^2$  by  $\log k$  gives the first-order answer  $d - n = rd$ , that is  $\text{ESS} \approx d(1 - r)$ .

*Step 5: the error term.* Let  $k = d - n$ . As  $rd \rightarrow \infty$ ,  $\psi(k/2) = \psi(d/2) + \log r \rightarrow \infty$  forces  $k \rightarrow \infty$ . The log of a  $\chi_k^2$  variable averages a little below the log of its mean:  $\psi(z) = \log z - \frac{1}{2z} + O(z^{-2})$  as  $z \rightarrow \infty$ . Hence

$$\log \frac{k}{d} - \frac{1}{k} + \frac{1}{d} + O(k^{-2}) = \log r \implies k = rd \left( 1 + \frac{1}{k} - \frac{1}{d} + O(k^{-2}) \right).$$

So  $k = rd(1 + o(1))$  and  $rd/k = 1 + O(1/k)$ . Multiplying out,  $k = rd + 1 - r + O(1/(rd))$ , and  $\text{ESS} = d - k = (d - 1)(1 - r) + O(1/(rd))$ . The correction  $-(1 - r)$  comes from the  $-1/k + 1/d$  term: the  $d - n$  untouched directions are fewer, so their chi-square sits further below its mean on the log scale.  $\square$

## B Proof of Proposition 2

*Idea.* Compare the learner with a twin who is also told whether the description is correct. The twin's regret is easy to write down. The learner cannot beat the twin, because extra information never hurts a Bayes-optimal predictor. And the learner cannot be much worse: its predictive puts weight  $p$  on the twin's answer when the description is correct and  $1 - p$  when it is wrong, so it loses at most  $\log(1/p)$  or  $\log(1/(1 - p))$  nats against the twin, which averages to  $H(p)$ .

*Proof.* Let  $Z = 1$  if the description is correct and  $Z = 0$  otherwise.

*Step 1: the predictive is a  $p$ -mixture.* Seeing  $m$  says nothing about  $Z$ , because  $m$  has the same law  $\mathcal{N}(0, (1-r)a_0I)$  either way, and  $x$  is independent of  $Z$ . With no examples there is nothing else to update on, so the posterior weight on  $Z = 1$  stays  $p$  and the Bayes predictive is

$$q(y) = p f_1(y) + (1-p) f_0(y), \quad f_1 = \mathcal{N}(m^\top x, 1 + r a_0 \|x\|^2), \quad f_0 = \mathcal{N}(0, 1 + a_0 \|x\|^2),$$

where  $f_1$  is the predictive given a correct description and  $f_0$  the predictive under the base prior.

*Step 2: the twin.* Given  $Z$ , the true conditional density of  $y$  is  $f_Z$ . So the twin predicts with  $f_1$  when  $Z = 1$ , with regret  $R_{p=1}^{\text{desc}}(r)$ , and with  $f_0$  when  $Z = 0$ , with regret  $R^{\text{ex}}(0)$ . Its regret is  $p R_{p=1}^{\text{desc}}(r) + (1-p) R^{\text{ex}}(0)$ .

*Step 3: the learner cannot beat the twin (lower bound).* Given  $(x, m, Z)$ , the expected log loss of any density  $q$  exceeds that of the true density  $f_Z$  by  $\text{KL}(f_Z \| q) \geq 0$ . Averaging over  $Z$  gives  $R^{\text{desc}}(p, r) \geq p R_{p=1}^{\text{desc}}(r) + (1-p) R^{\text{ex}}(0)$ .

*Step 4: the learner is at most  $H(p)$  worse (upper bound).* Pointwise,  $q \geq p f_1$  and  $q \geq (1-p) f_0$ . So when  $Z = 1$ ,  $-\log q \leq -\log f_1 + \log(1/p)$ , and when  $Z = 0$ ,  $-\log q \leq -\log f_0 + \log(1/(1-p))$ . Averaging, the learner's regret exceeds the twin's by at most  $p \log(1/p) + (1-p) \log(1/(1-p)) = H(p)$ .

*Step 5: the limits.*  $R_{p=1}^{\text{desc}}(r) = \frac{1}{2} \mathbb{E} \log(1 + r a_0 \|x\|^2)$  is nonnegative, which gives the floor  $(1-p) R^{\text{ex}}(0)$  for every  $r$ , and it decreases to 0 as  $r \rightarrow 0$  by monotone convergence, which gives the two limits.  $\square$